

ELL405 Operating Systems Assignment 2

Pratyush Pandey (2017EE10938)
Aniket Ulhasshetty (2018EE10443)

20nd March, 2021

Contents

1	Memory Management in XV6	1
1.1	Initializing the memory subsystem	1
1.2	Creating user processes	2
2	Implementation	2
2.1	fs.c	2
2.2	kalloc.c	4
2.3	proc.c	4
2.4	vm.c	4
2.5	Miscellaneous	8
3	myMemTest.c	8
3.1	Fork Test	8
3.2	Global protocols test	9
4	myMemTest1.c	10
4.1	SCFIFO test	10

1 Memory Management in XV6

1.1 Initializing the memory subsystem

The basic understanding of the way memory management is done in XV6 can be obtained using the following summary point

- xv6 uses 32-bit virtual addresses, resulting in a virtual address space of 4GB. xv6 uses paging to manage its memory allocations. However, xv6 does not do demand paging, so there is no concept of virtual memory.
- xv6 uses a page size of 4KB, and a two level page table structure.
- Remember that once the MMU is turned on, for any memory to be usable, the kernel needs a virtual address and a page table entry to refer to that memory location. When main starts, it is still using `entrypgdir` which only has page table mappings for the first 4MB of kernel. If the kernel wants to use more than this 4MB, it needs to map all of that memory as free pages into its address space, for which it needs a larger page table. So, main first creates some free pages in this 4MB in the function `kinit1`, which eventually calls the functions `freerange` and `kfree`.
- The kernel uses the struct `run` data structure to address a free page. This structure simply stores a pointer to the next free page, and the rest of the page is filled with garbage.

- The function `setupkvm` works as follows. For each of the virtual to physical address mappings in `kmap`, it calls `mappages`. The function `mappages` walks over the entire virtual address space in 4KB page-sized chunks, and for each such logical page, it locates the PTE using the `walkpgdir` function. `walkpgdir` simply outputs the translation that the MMU would do. It uses the first 10 bits to index into the page table directory to find the inner page table. If the inner page table does not exist, it requests the kernel for a free page, and initializes the inner page table. Note that the kernel has a small pool of free pages setup by `kinit1` in the first 4MB address space—these free pages are used to construct the kernel's page table. Once `walkpgdir` returns the PTE, `mappages` sets up the appropriate mapping using the physical address it has.

1.2 Creating user processes

- The function `userinit` creates the first user process. We will examine the memory management in this function. The kernel page table of this process is created using `setupkvm` as always. For the user part of the memory, the function `inituvm` (line 1903) allocates one physical page of memory, copies the `init` executable into that memory, and sets up a page table entry for the first page of the user virtual address space.
- If the child wants to execute a different executable from the parent, it calls `exec` right after `fork`. For example, the `init` process forks a child and execs the shell in the child. The `exec` system call copies the binary of an executable from the disk to memory and sets up the user part of the address space and its page tables.
- `Exec` first reads the ELF header of the executable from the disk and checks that it is well formed. It then initializes a page table, and sets up the kernel mappings in the new page table via a call to `setupkvm`. Then, it proceeds to build the user part of the memory image via calls to `allocuvm` and `loaduvm` for each segment of the binary executable. `allocuvm` allocates physical pages from the kernel's free pool via calls to `kalloc`, and sets up page table entries. `loaduvm` reads the memory executable from disk into the allotted page using the `readi` function.
- Next, `exec` goes on to build the rest of its new memory image. The guard page has no physical memory frame allocated to it, so any access beyond the stack into the guard page will cause a page fault. Then, the arguments to `exec` are pushed onto the user stack, so that the `exec` binary can access them when it starts.

2 Implementation

All the macros and small modifications have been added to the code, and described above, and can be seen using a git diff with the base code of XV6. However, the key/critical implementations have been written and explained. We make changes to the following files to add functionalities -

2.1 fs.c

- `readFromSwapFile`: Returns `sys_read` (-1 when error)
- `writeToSwapFile`: Returns `sys_write` (-1 when error)
- `getFreeSlot`: Returns free slot in file (-1 when full)
- `writePageToFile`: Place data to file
- `readPageFromFile`: Read data from file (-1 if wrong page address)
- `copySwapFile`: 1 if `swapFile` not in use, copy contents

```

1  /**
2  * remove swap file of proc p
3  * @param p
4  * @return
5  */
6  int removeSwapFile(struct proc* p){ ... }
7
8  /**
9  * return as sys_read (-1 when error)
10 * @param p
11 * @param buffer
12 * @param placeOnFile
13 * @param size
14 * @return
15 */
16 int readFromSwapFile(struct proc * p, char* buffer, uint placeOnFile, uint size){
17     ... }
18
19 /**
20 * return as sys_write (-1 when error)
21 * @param p
22 * @param buffer
23 * @param placeOnFile
24 * @param size
25 * @return
26 */
27 int writeToSwapFile(struct proc * p, char* buffer, uint placeOnFile, uint size){
28     p->swapFile->off = placeOnFile;
29     return filewrite(p->swapFile, buffer, size);
30 }
31
32 /**
33 *
34 * @param p
35 * @param userPageVAddr
36 * @param pgdir
37 * @return
38 */
39 int writePageToFile(struct proc * p, int userPageVAddr, pde_t *pgdir) {
40     int freePlace = getFreeSlot(p);
41     int retInt = writeToSwapFile(p, (char*)userPageVAddr, PGSIZE*freePlace, PGSIZE);
42     if (retInt == -1)
43         return -1;
44     p->fileCtrlr[freePlace].state = USED;
45     p->fileCtrlr[freePlace].userPageVAddr = userPageVAddr;
46     p->fileCtrlr[freePlace].pgdir = pgdir;
47     p->fileCtrlr[freePlace].accessCount = 0;
48     p->fileCtrlr[freePlace].loadOrder = 0;
49     return retInt;
50 }
51
52 /**
53 *
54 * @param p
55 * @param ramCtrlrIndex
56 * @param userPageVAddr
57 * @param buff
58 * @return
59 */
60 int readPageFromFile(struct proc * p, int ramCtrlrIndex, int userPageVAddr, char*
61     buff) { ... }
62
63 /**
64 *
65 * @param p
66 * @return 0 on success
67 */
68 int createSwapFile(struct proc* p){ ... }
69
70 /**

```

```

70 *
71 * @param fromP
72 * @param toP
73 */
74 void copySwapFile(struct proc* fromP, struct proc* toP){ ... }

```

2.2 kalloc.c

- `getFreePages`: Returns free pages
- `getTotalPages`: Returns total number of pages

2.3 proc.c

- `getPagedOutAmout`: flag indicating if page is swapped out
- `updateLap`: proc is either running, runnable or sleeping, update access counters
- `initSwapStructs`: Constructor/Initialiser

```

1 /**
2 *
3 * @param p
4 */
5 void initSwapStructs(struct proc* p) {
6     int i;
7     for (i = 0; i < MAX_TOTALPAGES - MAX_PYSC.PAGES; i++)
8         p->fileCtrlr[i].state = NOTUSED;
9 }
10 **
11 *
12 * @param p
13 * @return
14 */
15 int getPagedOutAmout(struct proc* p){
16
17     int i;
18     int amout = 0;
19
20     for (i=0; i < MAX_PYSC.PAGES; i++){
21         if (p->fileCtrlr[i].state == USED)
22             amout++;
23     }
24     return amout;
25 }
26
27 /**
28 *
29 */
30 void updateLap(){
31     struct proc *p;
32     acquire(&ptable.lock);
33     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
34         if (p->pid > 2 && p->state > 1 && p->state < 5) //proc is either running,
35             runnable or sleeping
36             updateAccessCounters(p); //implemented in vm.c
37     }
38     release(&ptable.lock);
39 }

```

2.4 vm.c

- `getPagePAddr`: flag indicating if page is swapped out
- `fixPagedOutPTE`: clear junk physical address, refresh CR3 register

- `fixPagedInPTE`: This method cannot be replaced with mappages because mappages cannot turn off `PTE_PG` bit
- `pageIsInFile`: PAGE IS IN FILE
- `getLIFO`
- `getSCFIFO`
- `getLAP`
- `getPageOutIndex`
- `updateAccessCounters`
- `getPageFromFile`
- `getFreeRamCtrlrIndex`
- `addToRamCtrlr`
- `swap`
- `isNONEpolicy`
- `removeFromRamCtrlr`
- Added a lot of additional code in other functions of the file `vm.c`

```

1  /**
2  *
3  * @param userPageVAddr
4  * @param pgdir
5  * @return
6  */
7  /**
8  int getPagePAddr(int userPageVAddr, pde_t * pgdir){
9      pte_t *pte;
10     pte = walkpgdir(pgdir, (int*)userPageVAddr, 0);
11     if(!pte) //uninitialized page table
12         return -1;
13     return PTEADDR(*pte);
14 }
15
16 /**
17 *
18 * @param userPageVAddr
19 * @param pgdir
20 */
21 void fixPagedOutPTE(int userPageVAddr, pde_t * pgdir){
22     ...
23 }
24
25 /**
26 * This method cannot be replaced with mappages because mappages cannot turn off
27 * PTE_PG bit
28 * @param userPageVAddr
29 * @param pagePAddr
30 * @param pgdir
31 */
32 void fixPagedInPTE(int userPageVAddr, int pagePAddr, pde_t * pgdir){
33     ...
34 }
35 /**
36 *
37 * @param userPageVAddr
38 * @param pgdir
39 * @return
40 */

```

```

41 int pageIsInFile(int userPageVAddr, pde_t * pgdir) {
42     pte_t *pte;
43     pte = walkpgdir(pgdir, (char *)userPageVAddr, 0);
44     return (*pte & PTE_PG); //PAGE IS IN FILE
45 }
46
47 /**
48  *
49  * @return
50  */
51 int getLIFO(){
52     int i;
53     int pageIndex = -1;
54     uint loadOrder = 0;
55
56     for (i = 0; i < MAX_PYSC_PAGES; i++) {
57         if (myproc()->ramCtrlr[i].state == USED && myproc()->ramCtrlr[i].loadOrder
58             > loadOrder) {
59             loadOrder = myproc()->ramCtrlr[i].loadOrder;
60             pageIndex = i;
61         }
62     }
63     return pageIndex;
64 }
65
66 /**
67  *
68  * @return
69  */
70 int getSCFIFO(){
71     ...
72 }
73
74 /**
75  *
76  * @return
77  */
78 int getLAP(){
79     int i;
80     int pageIndex = -1;
81     uint minAccess = 0xffffffff;
82
83     for (i = 0; i < MAX_PYSC_PAGES; i++) {
84         if (myproc()->ramCtrlr[i].state == USED && myproc()->ramCtrlr[i].
85             accessCount <= minAccess) {
86             minAccess = myproc()->ramCtrlr[i].accessCount;
87             pageIndex = i;
88         }
89     }
90     return pageIndex;
91 }
92
93 /**
94  *
95  * @return
96  */
97 int getPageOutIndex(){
98     #if LIFO
99         return getLIFO();
100     #endif
101     #if SCFIFO
102         return getSCFIFO();
103     #endif
104     #if LAP
105         return getLAP();
106     #endif
107     panic("Unrecognized paging mechanism");
108 }
109
110 /**
111  *
112  * @param p

```

```

111 */
112 void updateAccessCounters(struct proc * p){
113     ...
114 }
115
116 /**
117  *
118  * @return
119  */
120 int getFreeRamCtrlrIndex() {
121     if (myproc() == 0)
122         return -1;
123     int i;
124     for (i = 0; i < MAX_PYSC_PAGES; i++) {
125         if (myproc()->ramCtrlr[i].state == NOTUSED)
126             return i;
127     }
128     return -1; //NO ROOM IN RAMCTRLR
129 }
130
131 static char buff[PGSIZE]; //buffer used to store swapped page in getPageFromFile
132                               method
133
134 /**
135  *
136  * @param cr2
137  * @return
138  */
139 int getPageFromFile(int cr2){
140     ...
141 }
142
143 /**
144  *
145  * @param pgdir
146  * @param userPageVAddr
147  */
148 void addToRamCtrlr(pde_t *pgdir, uint userPageVAddr) {
149     ...
150 }
151
152 /**
153  *
154  * @param pgdir
155  * @param userPageVAddr
156  */
157 void swap(pde_t *pgdir, uint userPageVAddr){
158     ...
159 }
160
161 /**
162  *
163  * @return
164  */
165 int isNONEpolicy(){
166     #if NONE
167         return 1;
168     #endif
169     return 0;
170 }
171
172 /**
173  *
174  * @param userPageVAddr
175  * @param pgdir
176  */
177 void removeFromRamCtrlr(uint userPageVAddr, pde_t *pgdir){
178     ...
179 }
180
181 /**
182  *

```

```

182 * @param userPageVAddr
183 * @param pgdir
184 */
185 void removeFromFileCtrlr(uint userPageVAddr, pde_t *pgdir){
186 ...
187 }
188
189

```

2.5 Miscellenous

- The functions `isdireempty` and `create` in `sysfile.c` were made non static and added to `defs.h` for access to members in other files
- Added Page Fault conditions in `trap.c`
- Added `myMemTest.c` in Makefile
- Added `myMemTest1.c` with additional tests
- Added additional selection macros in Makefile
- Added definitions in `proc.h`
- Added bits flags in `mmu.h`

3 myMemTest.c

3.1 Fork Test

```

1  *
2  Test used to check the swapping machanism in fork.
3  Best tested when LIFO is used (for more swaps)
4  */
5  void forkTest(){
6      int i;
7      char * arr;
8      arr = malloc(50000); //allocates 13 pages (sums to 16), in lifo , OS puts page
9                             #15 in file.
10
11      for (i = 0; i < 50; i++) {
12          arr[49100+i] = 'A'; //last six A's stored in page #16, the rest in #15
13          arr[45200+i] = 'B'; //all B's are stored in page #15.
14      }
15      arr[49100+i] = 0; //for null terminating string...
16      arr[45200+i] = 0;
17
18      if (fork() == 0){ //is son
19          for (i = 40; i < 50; i++) {
20              arr[49100+i] = 'C'; //changes last ten A's to C
21              arr[45200+i] = 'D'; //changes last ten B's to D
22          }
23          printf(1, "SON: %s\n",&arr[49100]); // should print AAAAA..CCC...
24          printf(1, "SON: %s\n",&arr[45200]); // should print BBBBB..DDD...
25          printf(1, "\n");
26          free(arr);
27          exit();
28      } else { //is parent
29          wait();
30          printf(1, "PARENT: %s\n",&arr[49100]); // should print AAAAA...
31          printf(1, "PARENT: %s\n",&arr[45200]); // should print BBBBB...
32          free(arr);
33      }
34  }
35  /**
36  * The Output of the above test is as follows:
37  SON: AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA

```



```

37 SON: BBBB...
38
39 PARENT: AAAAA...
40 PARENT: BBBB...
41
42 * Hence the test is successful
43 */
44

```

```

xv6...
cpu1: starting
cpu0: starting
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ myMemTest
SON: AAAAA...
SON: BBBB...

PARENT: AAAAA...
PARENT: BBBB...
$ █

```

Figure 1: myMemTest output

3.2 Global protocols test

```

1  /*
2  Global Test:
3  Allocates 17 pages (1 code, 1 space, 1 stack, 14 malloc)
4  Using pseudoRNG to access a single cell in the array and put a number in it.
5  Idea behind the algorithm:
6  Space page will be swapped out sooner or later with scfifo or lap.
7  Since no one calls the space page, an extra page is needed to play with swapping
   (hence the #17).
8  We selected a single page and reduced its page calls to see if scfifo and lap
   will become more efficient.
9  */
10 void globalTest() {
11     char * arr;
12     int i;
13     int randNum;
14     arr = malloc(ARR_SIZE); //allocates 14 pages (sums to 17 - to allow more then one
   swapping in scfifo)
15     for (i = 0; i < TEST_POOL; i++) {
16         randNum = getRandNum(); //generates a pseudo random number between 0 and
   ARR_SIZE
17         while (PGSIZE*10-8 < randNum && randNum < PGSIZE*10+PGSIZE/2-8)
18             randNum = getRandNum(); //gives page #13 50% less chance of being selected
   //(redraw number if randNum is in the first half of
19             page #13)
20         arr[randNum] = 'X'; //write to memory
21     }
22     free(arr);
23 }
24 /**
25 * The Output of the above test is as follows:
26 Results (for TEST_POOL = 500):
27 LIFO: 42 Page faults
28 LAP: 18 Page faults
29 SCFIFO: 35 Page faults
30 * Hence the test is successful
31 */
32

```

4 myMemTest1.c

Here we have additional tests

4.1 SCFIFO test

```

1  int i, j;
2  char *arr[14];
3  char input[10];
4
5  // TODO delete
6  printf(1, "myMemTest: testing SCFIFO... \n");
7
8  // Allocate all remaining 12 physical pages
9  for (i = 0; i < 12; ++i) {
10     arr[i] = sbrk(PGSIZE);
11     printf(1, "arr[%d]=0x%x\n", i, arr[i]);
12 }
13 printf(1, "Called sbrk(PGSIZE) 12 times — all physical pages taken.\nPress any
    key...\n");
14 gets(input, 10);
15
16 /*
17 Allocate page 15.
18 For this allocation, SCFIFO will consider moving page 0 to disk, but because it
19 has been accessed, page 1 will be moved instead.
20 Afterwards, page 1 is in the swap file, the rest are in memory.
21 */
22 arr[12] = sbrk(PGSIZE);
23 printf(1, "arr[12]=0x%x\n", arr[12]);
24 printf(1, "Called sbrk(PGSIZE) for the 13th time, no page fault should occur and
    one page in swap file.\nPress any key...\n");
25 gets(input, 10);
26
27 /*
28 Allocate page 16.
29 For this allocation, SCFIFO will consider moving page 2 to disk, but because it
30 has been accessed, page 3 will be moved instead.
31 Afterwards, pages 1 & 3 are in the swap file, the rest are in memory.
32 */
33 arr[13] = sbrk(PGSIZE);
34 printf(1, "arr[13]=0x%x\n", arr[13]);
35 printf(1, "Called sbrk(PGSIZE) for the 14th time, no page fault should occur and
    two pages in swap file.\nPress any key...\n");
36 gets(input, 10);
37
38 /*
39 Access page 3, causing a PGFLT, since it is in the swap file. It would be
40 hot-swapped with page 4. Page 4 is accessed next, so another PGFLT is invoked,
41 and this process repeats a total of 5 times.
42 */
43 for (i = 0; i < 5; i++) {
44     for (j = 0; j < PGSIZE; j++)
45         arr[i][j] = 'k';
46 }
47 printf(1, "5 page faults should have occurred.\nPress any key...\n");
48 gets(input, 10);
49
50 /*
51 If DEBUG flag is defined as != 0 this is just another example showing
52 that because SCFIFO doesn't page out accessed pages, no needless page faults
53 occur.
54 */
55 if(DEBUG){
56     for (i = 0; i < 5; i++) {
57         printf(1, "Writing to address 0x%x\n", arr[i]);
58         arr[i][0] = 'k';
59     }
60 }
61 //printf(1, "No page faults should have occurred.\nPress any key...\n");
62 gets(input, 10);

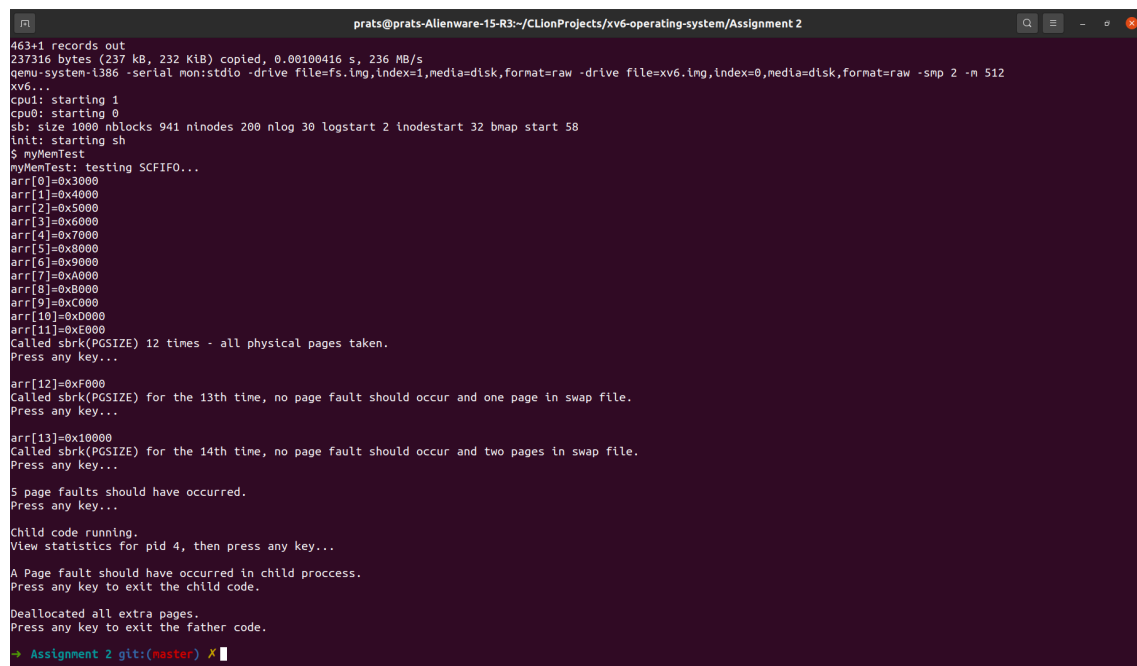
```

```

59 }
60
61 if (fork() == 0) {
62     printf(1, "Child code running.\n");
63     printf(1, "View statistics for pid %d, then press any key...\n", getpid());
64     gets(input, 10);
65
66     /*
67     The purpose of this write is to create a PGFLT in the child process, and
68     verify that it is caught and handled properly.
69     */
70     arr[5][0] = 'k';
71     printf(1, "A Page fault should have occurred in child process.\nPress any key
72     to exit the child code.\n");
73     gets(input, 10);
74
75     exit();
76 }
77 else {
78     wait();
79
80     /*
81     Deallocate all the pages.
82     */
83     sbrk(-14 * PGSIZE);
84     printf(1, "Deallocated all extra pages.\nPress any key to exit the father code
85     .\n");
86     gets(input, 10);
87 }

```

The result is shown below



```

prats@prats-Alienware-15-R3:~/CLionProjects/xv6-operating-system/Assignment 2
463+1 records out
237316 bytes (237 kB, 232 KiB) copied, 0.00100416 s, 236 MB/s
qemu-system-i386 -serial mon:stdio -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,index=0,media=disk,format=raw -smp 2 -n 512
xv6...
cpu0: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ myMemTest
myMemTest: testing SCFIFO...
arr[0]=0x3000
arr[1]=0x4000
arr[2]=0x5000
arr[3]=0x6000
arr[4]=0x7000
arr[5]=0x8000
arr[6]=0x9000
arr[7]=0xA000
arr[8]=0xB000
arr[9]=0xC000
arr[10]=0xD000
arr[11]=0xE000
Called sbrk(PGSIZE) 12 times - all physical pages taken.
Press any key...
arr[12]=0xF000
Called sbrk(PGSIZE) for the 13th time, no page fault should occur and one page in swap file.
Press any key...
arr[13]=0x10000
Called sbrk(PGSIZE) for the 14th time, no page fault should occur and two pages in swap file.
Press any key...
5 page faults should have occurred.
Press any key...
Child code running.
View statistics for pid 4, then press any key...
A Page fault should have occurred in child process.
Press any key to exit the child code.
Deallocated all extra pages.
Press any key to exit the father code.
+ Assignment 2 git:(master) X

```

Figure 2: Extra SCFIFO test output